Chapter10 Tips for Deep Learning

1. Bad Results on Training Data?
   a) New activation function
      i) Vanishing Gradient Problem

         In the layers near input, gradients are small, learn very slow, and the parameters are almost random, whereas in the layers near output, gradients are large, learn very fast, and the parameters are almost converged

         Intuitive reason: Sigmoid function: large input, small output

      ii) Alternative method:

         ReLU:

         x<0: y = 0; x>0: y = x

         1. Fast to compute
         2. Biological reason
         3. Handle the problem of vanishing gradient

         Leaky ReLU:

         x<0: y = 0.01x; x>0: y = x

         Parametric ReLU:

         x<0: y = αx; x>0: y = x

         α is learned by gradient descent

         Maxout:

         1. ReLU is a special case of Maxout
         2. Learnable activation function
            a) Activation function in Maxout network can be any piecewise linear convex function
            b) How many pieces depending on how many elements in a group

   b) Adaptive learning rate
      i) Adagrad

      1) $w^{t+1} = w^t - \frac{\eta}{\sqrt{\sum_{i=0}^{t}(g^i)^2}} g^t$

      2) Use first derivative to estimate second derivative

      ii) RMSProp

      1) Error surface can be very complex when training DNN

      2) $w^1 = w^0 - \frac{\eta}{\sigma^0} g^0 \quad \sigma^0 = g^0$

      $w^2 = w^1 - \frac{\eta}{\sigma^1} g^1 \quad \sigma^1 = \sqrt{\alpha(\sigma^0)^2 + (1-\alpha)(g^1)^2}$

      $w^{t+1} = w^t - \frac{\eta}{\sigma^t} g^t \quad \sigma^t = \sqrt{\alpha(\sigma^{t-1})^2 + (1-\alpha)(g^t)^2}$

iii) Momentum

1) Movement not just based on gradient, but previous movement
2) Movement = momentum + negative of $\partial L / \partial w$

iv) Adam

RMSProp + Momentum

2.Bad Results on Testing Data?

a) Early Stopping
b) Regularization
   i) L1 Regularization
      1) $L'(\theta) = L(\theta) + \lambda||\theta||_1 \qquad ||\theta||_1 = |w_1| + |w_2| + \cdots$
      2) $\frac{\partial L'(\theta)}{\partial w} = \frac{\partial L(\theta)}{\partial w} + \lambda sgn(w)$
      3) $w^{t+1} = w^t - \eta \frac{\partial L}{\partial w} - \eta\lambda sgn(w^t)$
   ii) L2 Regularization
      1) $L'(\theta) = L(\theta) + \frac{1}{2}\lambda||\theta||_2 \qquad ||\theta||_2 = (w_1)^2 + (w_2)^2 + \cdots$
      2) $\frac{\partial L'(\theta)}{\partial w} = \frac{\partial L(\theta)}{\partial w} + \lambda w$
      3) $w^{t+1} = (1 - \eta\lambda)w^t - \eta \frac{\partial L}{\partial w}$
   iii) Weight Decay
c) Dropout
   i) Training:
      1) Each time before updating the parameters, each neuron has p% to dropout, the structure of the network is changed
      2) Using the new network for training
   ii) Testing:
      1) No Dropout
      2) If the dropout rate at training is p%, all the weights times 1-p%
   iii) Dropout is a kind of ensemble
      1) Train a bunch of networks with different structures
      2) Using one mini-batch to train one network
      3) Some parameters are shared in the network